



If one puts an infinite number of monkeys in front of (strongly built) typewriters, and lets them clap away, there is a certainty that one of them would come out with an exact version of the *Iliad*. Upon examination, this may be less interesting a concept than it appears at first: such probability is very low. But let us carry the reasoning one step beyond. Now that we have found that hero among monkeys, would any reader invest his life's savings on a bet that the monkey would write the *Odyssey* next?

In this story, it is the second step that is interesting. How much can past performance (here the typing of the *Iliad*) be relevant in forecasting future performance? The same applies to any decision based on past performance, merely relying on the attributes of the past time series. Think about the monkey showing up at your door with his impressive past performance. Hey, he wrote the *Iliad*. Quickly, sign him up for the sequel.

The major problem with inference in general is that those whose profession is to derive conclusions from data often fall into the trap faster and more confidently than others. The more data we have, the more likely we are to drown in it. For common wisdom among people with a budding knowledge of probability laws is to base their decision-making on the following principle: it is very unlikely for someone to