



Resource Utilisation

Computationally intensive during the training phase, requiring significant processing power and time.

Requires more resources during inference for the retrieval step, potentially increasing latency.



Domain Adaptation

Excellent adaptation to specialised domains, producing highly relevant outputs for specific fields.

Can dynamically adapt to a variety of domains by utilising up-to-date external knowledge bases without retraining.



Maintenance

Updates necessitate retraining the model with new data to incorporate changes and improvements.

It is more straightforward to update by modifying or expanding the retrieval database with new information.



Use Cases

Ideal for specialised NLP tasks such as legal document analysis, medical report generation, and domain-specific chatbots.

Suitable for open-domain question answering, providing current information, and dynamic content generation.