

models that did best in the out-of-sample period. It is one of the most common errors people make and one of the reasons why data mining as it is typically applied yields terrible results.

What should you be doing instead?

You can look for patterns where, on average, all the models out-of-sample continue to do well. You know you are doing well if the average for the out-of-sample models is a significant percentage of the in-sample score. Generally speaking, you are really getting somewhere if the out-of-sample results are more than 50 percent of the in-sample. QIM's business model would never have worked if SAS and IBM were building great predictive modeling software.

Because if they did, lots of people could use this software for financial modeling?

And lots of people do, but they have a very tough time modeling properly with the software, and they end up doing data mining of the bad kind.

Why do you think you as a single individual have been able to come up with data mining procedures that are much more effective in financial markets than the software produced by these large corporations with tons of PhDs?

Because the commercial software is focused more on the problem of allowing users to handle large quantities of data than on providing users with very strict protocols to make sure they don't curve-fit the data. People are so excited about building and using software that allows them to handle so much more data than ever before that they are missing the point of doing the process properly. Not only does the software fail to guide the user in doing data mining correctly, it actually leads users in the wrong direction because it allows them to generate bogus evidence to support their pet theories.

Do you give the same weight to data from the 1980s as data from the 2000s?

Sometimes we give a little more weight to more recent data, but it is amazing how valuable older data still is. The stationarity of the patterns