



## Behaviour Evaluation and Assessment

As IBM continues to push the boundaries of generative AI, the need for robust evaluation and assessment mechanisms becomes increasingly critical. To address this challenge, IBM is developing a comprehensive suite of harm and hallucination detectors and leveraging the power of large language models (LLMs) as judges.

These cutting-edge tools are designed to identify and mitigate potential risks associated with generative AI outputs, such as the generation of harmful or biased content or the propagation of misinformation. By incorporating these safeguards into its platform, IBM aims to instill a heightened level of trust and confidence in the generative AI solutions it delivers to its clients.

Furthermore, IBM is pioneering innovative source attribution and red-team audit capabilities, enabling organisations to trace the origins of generative AI outputs and validate their authenticity. These capabilities are poised to be delivered to Watsonx and the data & model factory, further solidifying IBM's position as a responsible and transparent AI development leader.